

Understanding Stragglers in Large Model Training Using What-if Analysis

Using a 5-Month Production Trace from ByteDance

Jinkun Lin, New York University; Ziheng Jiang, Zuquan Song, Sida Zhao, and Menghan Yu, ByteDance Seed; Zhanghan Wang, New York University; Chenyuan Wang, ByteDance Seed; Zuocheng Shi, Zhejiang University; Xiang Shi, ByteDance; Wei Jia, Zherui Liu, Shuguang Wang, Haibin Lin, and Xin Liu, ByteDance Seed; Aurojit Panda and Jinyang Li, New York University


Anant Goyal

What Are Stragglers?

Problem: Distributed LLM training requires frequent synchronization and tighter coordination

- **Straggler:** Worker that lags behind peers, stalling the entire training job
- **Worker:** Single GPU + controlling CPU process

Why it matters: All workers must synchronize → slowest worker blocks everyone



Source: Google Cloud Blog

Scale of the Problem

- Empirically shown: larger model size = improved task accuracy
 - Meta Llama: 65B params (Llama-1) → 405B params (Llama-3)
- Traditional straggler mitigation **doesn't work**:
 - Backup workers → bad training performance (assumes infrequent synchronization)
 - Drop slow updates → hurts model accuracy

Key Question: Do stragglers pose performance issues in real-world LLM training?

Answer: YES!

Key Takeaways

 42.5% of jobs are $\geq 10\%$ slower due to stragglers; tail jobs waste **45% of GPU hours**

- Root causes are **software & workload**, not hardware failures (only 1.7% of cases)

 **Most Surprising:**

- **Compute**, not communication, is the bottleneck — especially on a well-optimized network
- Straggling has **no correlation with job size** — more GPUs doesn't mean more stragglers
- A simple **Python Garbage Collection tweak** was hiding in plain sight

Study Overview

Setup

- 3,079 training jobs with each using between 128 and > 5000 GPUs
- Network tuned to ensure **no slowdown from congestion**

Design Choice: Jobs don't share servers

- Eliminates stragglers from resource contention
- Any observed stragglers are due to training dynamics, **not infrastructure**

Estimating "Straggler-Free" Performance

Key Insight: Comparable operations should have same duration in absence of stragglers

1. Log (trace) operation type, start/end timestamps, metadata
2. Calculate **idealized duration** for a job, T_{ideal}
3. Calculate **slowdown** metrics and **GPU hours wasted**

$$S = \frac{T}{T_{ideal}} \quad \text{and} \quad \% \text{ resource waste} = 1 - \frac{1}{S}$$


How do we calculate T_{ideal} ?

For Compute Operations: Use **average** duration across all workers

- Rationale: Same workload → should take same time

For Communication Operations: Use **median** duration

- Rationale: Affected operations take very long (switch/NIC flapping)




Source: Original author slides from OSDI '25

Finding #1 - GPU Hours Wasted

- Stragglers cause **significant resource waste**
 - >50% of jobs: waste 7.8% of GPU hours
 - ~1% of jobs: waste 45% of GPU hours

At distributed LLM scale with thousands of GPUs:

- 7.8% waste = hundreds of GPU-hours
- GPU-hours wasted → tens of thousands \$



Finding #2 - Compute vs Communication


→ Most waste caused by **compute operations, NOT communication**

Why?

- ByteDance optimized in-house network
- Ample network bandwidth provisioned
- Network tuned to avoid congestion

⚠ Counterintuitive: Most assume network is the bottleneck

Cumulative distribution function (CDF) of resource wasted by operation type



Other Findings

Step-Level Slowdown is Persistent

- 90th percentile of steps: slowdown of only 1.06x
- Small per-step slowdowns compound over thousands of training steps

⚠ Sufficient to sample just a few steps to identify stragglers

No Correlation Between Job Size and Straggling ⚠ Counterintuitive!

- Larger jobs = more GPUs = *should* mean more straggler risk
- But data shows small and large jobs are equally affected
- Straggling is inherent to **training dynamics** (model type, workload), not scale

Root Cause #1 - Stage Partitioning Imbalance


Problem: Pipeline stages divided evenly by layer count

- Last pipeline stage includes **loss layer**
- Loss layer computation $>>$ transformer layers
- Last stage becomes bottleneck

Attempted Fix: Assign fewer layers to last stage

- Result: **<10% speedup** (minimal improvement) due to manual tuning

Cumulative distribution function (CDF) of slowdown caused by the last pipeline stage



Root Cause #2 - Python Garbage Collection


Problem: Forward compute operations run in Python with automatic GC

- "Stop-the-world" GC pauses Python code execution
- Different processes can trigger GC at different times

Fix: Turn off Python automatic GC

- Manually run GC every 500 training steps
- Result: **12.6% improvement**

Implemented, but not widely adopted due to memory challenges



Other Root Causes

Sequence Length Imbalance

- Microbatch compute time \propto sum of sequence length squares
- **21.4% of jobs** slowed down by sequence length imbalance
- Greedy balancing fix → **23.9% improvement** not deployed; doesn't fix PP-level imbalance

Uncommon Causes


- PyTorch CUDA memory allocator slowness → slows forward & backward pass
 - Unrelated kernels sharing CUDA hardware queue → false dependencies block kernel launch
- These are rare but can cause significant slowdowns in specific jobs

SMon - Straggler Monitoring System


Displays:

- Estimated overall slowdown per job
- Per-step slowdown trends
- Per-worker slowdown identification


Used in production by on-call team to **diagnose and identify stragglers in real-time**



(a) Worker issue



(b) Stage partitioning imbalance



(c) Sequence length imbalance

Related Work

Traditional Straggler Mitigation (MapReduce era)

- Mantri [OSDI '10], Dolly – backup tasks, speculative execution
 - ✗ Assumes infrequent sync; breaks under LLM's tight coordination requirements

Stragglers in LLM Training

- FALCON [arXiv 2024]
 - Shared cluster only → sees resource contention stragglers
 - Manual root cause analysis
 - Misses stragglers that affect **most steps**

This paper's contribution: A large-scale, semi-automated empirical study of stragglers in LLM training, with a deployed monitoring tool (SMon)

Study Limitations

Data Coverage

- **Discarded jobs:** repeated failures, trace issues, simulation discrepancy >5%
- May underrepresent worst-case stragglers and hardware-related failures

Profiling Tool (NDTimeline)

- Cannot capture stragglers **within TP or CP groups** — these slow all microbatches uniformly

Cluster-Specific Findings

- Network is **deliberately overprovisioned** — communication bottlenecks may matter more elsewhere

Discussion & Critique

Strengths:

- Large-scale empirical study (5 months, production workload)
- Counterfactual, What-if analysis is rigorous and novel
- Led to practical tool (SMon) and deployed fixes (Planned GC)

Questions

- Forward compute runs in Python, backward runs in C++ — **why not rewrite forward in C++?**
- Can stage partitioning be done **automatically** based on FLOPs rather than manual tuning?
- SMon detects stragglers — can it **actively mitigate** them (e.g., rebalance sequences at runtime)?
- How do findings change in **heterogeneous or shared clusters** where network is not overprovisioned?

Q: Why Not Deploy Sequence Length Balancing?

Result: Greedy balancing → 23.9% throughput improvement in experiments

Why not deployed?

- Only balances across **DP microbatches** — does not fix imbalance at the **PP level**
- Balanced sequence lengths ≠ balanced memory consumption
 - Long sequences require more activation memory → risk of OOM on some workers
- WLB-LLM (the companion paper) addresses this more completely

Takeaway: The fix works locally but doesn't solve the full problem

Q: How Does This Compare to FALCON?

| | This Paper | FALCON |
|-------------------------------------|---------------------------------------|------------|
| Cluster type | Dedicated | Shared |
| Resource contention stragglers | Not observed | Present |
| Large job traces (≥ 512 GPUs) | 562 jobs | 27 jobs |
| Analysis method | Semi-automated (simulation + what-if) | Manual |
| Persistent step stragglers | Studied | Overlooked |

Key insight: Shared vs. dedicated cluster changes which root causes dominate

- FALCON sees contention stragglers; this paper sees workload imbalance
- Neither result is "wrong" — they reflect different real-world deployment settings

WLB-LLM: Workload-Balanced 4D Parallelism for Large Language Model Training

*Zheng Wang, Anna Cai, Xinfeng Xie, Zaifeng Pan, Yue Guan, Weiwei Chu, Jie Wang,
Shikai Li, Jianyu Huang, Chris Cai, Yuchen Hao, Yufei Ding*

Kartik Ramesh


Introduction

- Llama 3.1 405B trained over 16k GPUs for 30M H100 hours.
- Using AWS H100 pricing **\$212M**.
- A 1.2x increase in training speed would have saved **\$36M** in cloud costs.
- WLLB-LLM is a work from Meta and UCSD, released ~6 months after their training Llama 3 that achieves this.

Self Attention

- Attention scales quadratically with the number of past tokens $O(T^2)$
- Not all tokens are equal: Processing later tokens in a document require more computation than earlier.
- When packing multiple documents for training, control attention span using masks.

Packed Docs Attention Mask



4D Parallelism

- Huge models are trained on huge clusters because of scale and necessity.
- Tensor shape $[B, T, H]$ across several layers.
- 4D Parallelism splits this shape in various ways.
 - Data - Split B across replicated model.
 - Pipeline - Split model into groups of layers.
 - Context - Split along sequence length T
 - Tensor - Split across the H .

4D Parallelism



$DP > PP > CP > TP$

Key Takeaways


Problem: In long-context 4D training, token count is a weak proxy for compute; attention cost is highly non-uniform.

- **Core idea #1 (PP):** Reduce PP imbalance via attention-aware micro-batch packing.
- **Core idea #2 (CP):** Reduce CP imbalance via fine-grained per-document sharding.

Bottom line: WLB-LLM improves training throughput ($\approx 1.23x$) without hurting convergence.


Motivation #1 - Pipeline Imbalance

- Pipeline Parallelism splits B into $N \mu\text{B}$ to hide delays.
- Balanced Batches \rightarrow Higher throughput




Motivation #1 - Pipeline Imbalance

- Naive approach → split by tokens.
- Each μB has uniform sequence length = CONTEXT_WINDOW_SIZE
- Does this balance LLM workload?
- $k.128^2 \gg k.128 * 1^2$



Motivation #2 - Context Imbalance

- Workload across CP workers should be balanced.
- After packing, split the sequence into $2 \cdot CP$ parts and assign one from front and one from back to balance attention.
- Good heuristic, fails for multiple packed documents. Common in long context training.
- Every small delay adds up to higher-order delays.



Baseline: Attention-Aware packing.

Idea: Divide B into μB by estimating d_i^2 as attention cost for each document.

- It works, but limited balancing improvements
→ limited speedup.
- Higher balancing across μB requires balancing across multiple global B . This disturbs the random order of training and loss convergence.
- It might be impossible to come up with such a μB construction, if there are no candidates.

$$\begin{aligned} & \text{minimize} && \max\left(\sum_{i=1}^N x_{ij} \cdot d_i^2\right), \quad j = 1, \dots, M \\ & \text{subject to} && \sum_{j=1}^M x_{ij} = 1, \quad i = 1, \dots, N \\ & && \sum_{i=1}^N x_{ij} \cdot d_i \leq L, \quad j = 1, \dots, M \\ & && x_{ij} \in \{0, 1\} \end{aligned}$$


Attention Cost
Each Document is present in only one batch
L = Context Size

Variable-Length Packing

Idea: Allow $\text{len}(\mu B) > \text{CONTEXT_SIZE}$ for weaker μB



- Attention is Quadratic, but other operations are linear (feed-forward, comms etc.)
- Balance the total workload, not just attention.
- Balance long documents against many shorter documents.

$$\min \left(\max \left(\sum_{i=1}^N \left(W_a(x_{ij} \cdot d_i) + W_l(x_{ij} \cdot d_i) \right) \right) \right)$$



Outlier document detection


- Still, you might not have sufficient smaller documents to balance the load of a long document.
- Observe: there aren't that many ultra-long documents.
- Instead of balancing across multiple batches, delay the few long documents.
- Model convergence should not hurt significantly.



Improved CP sharding

Idea: Apply CP indexing logic to each individual document.

- This should yield a more balanced workload across multiple CP workers.
- They also implement an optimization to avoid padding tokens.




Kernel inefficiencies

- Per-Document sharding achieves better balance, but it does not always guarantee better performance.
- Smaller per-rank attention problems reduce kernel efficiency:
 - Poor tile utilization → padding overhead for short sequences (<128 tokens).
 - Lower effective FLOPs utilization → higher time per token.
 - Reduced KV tile reuse → weaker Hopper TMA multicast benefits

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right) V$$

Kernel Inefficiencies



Experimental Setup

- **Cluster:** 32 nodes, each with 8x NVIDIA H100 SXM 80GB GPUs.
- **Interconnect:** NVLink intra-node, RoCE inter-node.
- **Models:** LLaMA-like 550M, 7B, 30B, 70B; each tested at 64K and 128K context.
- **Training config:** 4D parallelism, global batch size = `PP_size x DP_size`, bfloat16 precision.
- **Baselines:**
 - `Plain-4D` : default 4D training with per-sequence CP sharding.
 - `Fixed-4D` : fixed-length packing + fixed CP sharding (per-sequence or per-document).

Speedup Breakdown

Which optimization helps us the most?

- PP-Var-Len alone → 1.28x
- Orthogonal optimizations that combine well.
- Every second counts!




Figure 13: Performance breakdown of *WLB-LLM* on the 7B model with a 128K context window.

Speedup across Model + Context




Figure 12: Training performance speedups of *WLB-LLM* and *Fixed-4D* over *Plain-4D* across various configurations.

- WLB-LLM consistently outperforms baseline for all tested configurations.
- Naive attention balancing is insufficient.
- Relative speedup decreases with increased model size.

Other Experiments (Summary)

- **Context sensitivity (Fig. 14):** Speedup increases with context length (about $1.07x$ @64K to $1.40x$ @160K on 7B), consistent with worse imbalance at longer contexts.
- **Packing overhead vs balance (Table 2):** WLB-LLM reaches near-optimal imbalance with low runtime overhead (tens of ms), while solver-based packing can be prohibitively slow.
- **CP sharding ablation (Fig. 15):** Adaptive sharding consistently outperforms always-per-sequence or always-per-document sharding.
- **Convergence / quality:** Their training-loss curves indicate no clear quality regression from the system optimizations.

Discussion & Critique

Strengths

- Identifies and fixes a bottleneck for training long-context Llama models with 4D parallelism.
- Joint PP+CP optimization gives meaningful real-system gains, especially as context windows grow.
- Engineering is practical: low overhead and no obvious convergence regression in their reported runs.

Weaknesses

- Workload dependence and limited generalization evidence outside their evaluated distribution.
- Strong dependence on a small number of extreme outliers; unclear benefit when length distributions are flatter.
- Heavy use of heuristics (packing + sharding selection) without strong guarantees in worst-case settings.

Related Work

(1) Efficient Long-context Language Model Training by Core Attention Disaggregation (DistCA)

- Split out “core attention” ($\text{softmax}(QK^\top)V$) as a weightless compute service, separate from the rest of the transformer.
- Better than WLB-style baselines at scale: reports $\sim 1.15\text{--}1.35\times$ throughput gains over their WLB “ideal” baseline in 4D (with PP), depending on workload.

(2) ByteScale Efficient Scaling of LLM Training with a 2048K Context Length on More Than 12,000 GPUs

- Hybrid Data Parallelism (HDP): unify DP + CP into one dynamic device mesh.
- Length-aware sharding: use the minimum number of devices per sequence.
- Short sequences stay local (skip CP comm), long sequences shard across more GPUs.

(3) Ordering efficiency

- Reduce pipeline bubbles by optimizing scheduling.
- PipeDream, 1F1B, Seq1F1B.

What did you think?

Problem: In long-context 4D training, token count is a weak proxy for compute; attention cost is highly non-uniform.

- **Core idea #1 (PP):** Reduce PP imbalance via attention-aware micro-batch packing.
- **Core idea #2 (CP):** Reduce CP imbalance via fine-grained per-document sharding.

Bottom line: WLB-LLM improves training throughput ($\approx 1.23x$) without hurting convergence.

| Packing Method | | Imbalance Degree | Packing Overhead (ms) |
|-------------------------|-----------------|------------------|-----------------------|
| Method | Config | | |
| <i>Original Packing</i> | / | 1.44 | 0 |
| <i>Fixed-Len Greedy</i> | #global batch=1 | 1.41 | 4 |
| | #global batch=2 | 1.22 | 5 |
| | #global batch=4 | 1.11 | 5 |
| | #global batch=8 | 1.08 | 5 |
| <i>Fixed-Len Solver</i> | #global batch=1 | 1.40 | 467 |
| | #global batch=2 | 1.18 | 1488 |
| | #global batch=4 | 1.09 | 25313 |
| <i>WLB-LLM</i> | #queue=1 | 1.24 | 8 |
| | #queue=2 | 1.05 | 20 |
| | #queue=3 | 1.05 | 23 |

Table 2: Packing imbalance degree and overhead analysis.