

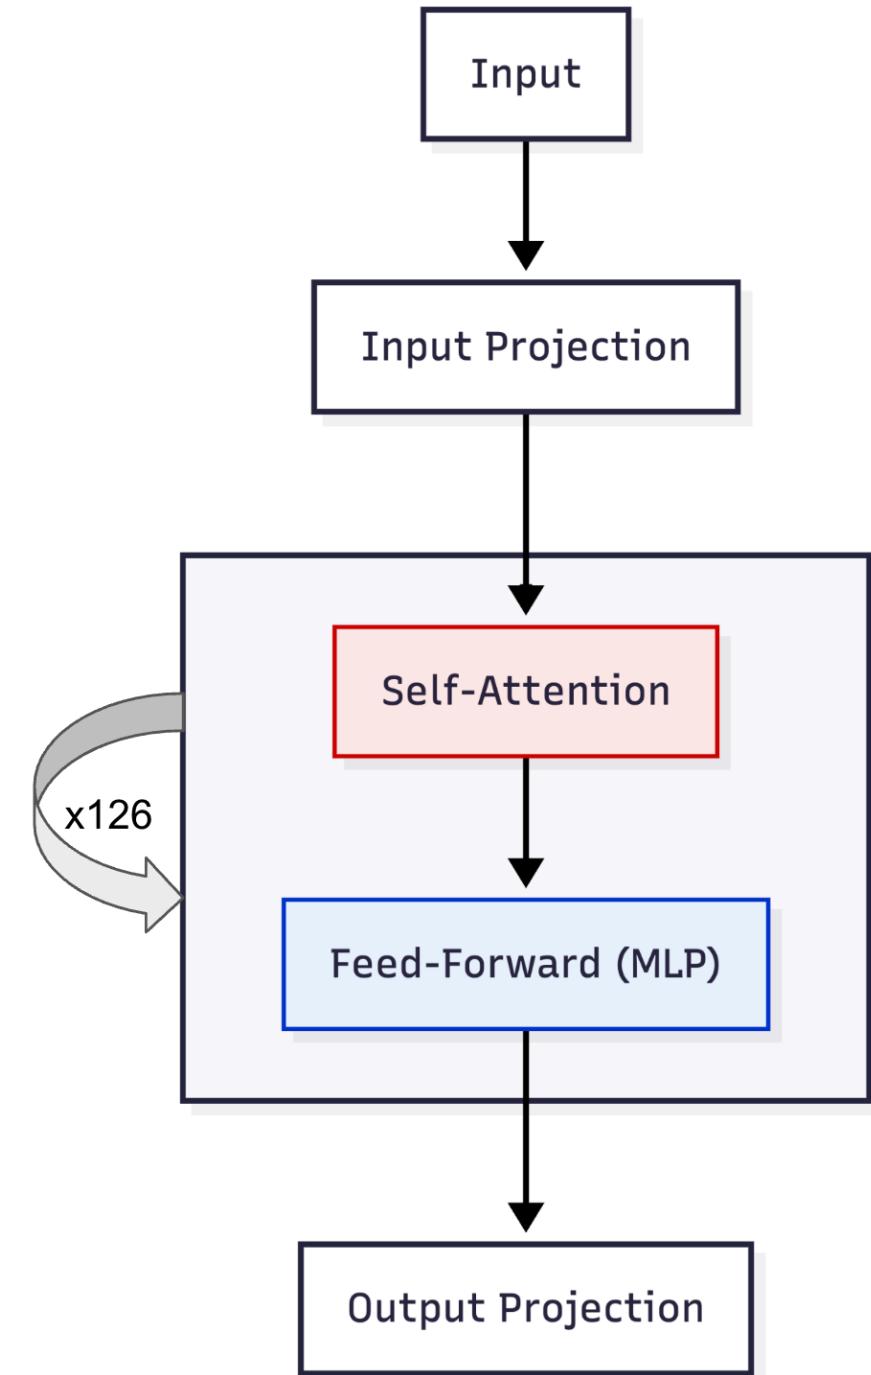
# Introduction

- Llama 3.1 405B trained over 16k GPUs for 30M H100 hours.
- Using AWS H100 pricing **\$212M**.
- A 1.2x increase in training speed would have saved **\$36M** in cloud costs.
- WLLM is a work from Meta and UCSD, released ~6 months after their training Llama 3 that achieves this.

# Background - Model Architecture

Things to keep in mind:

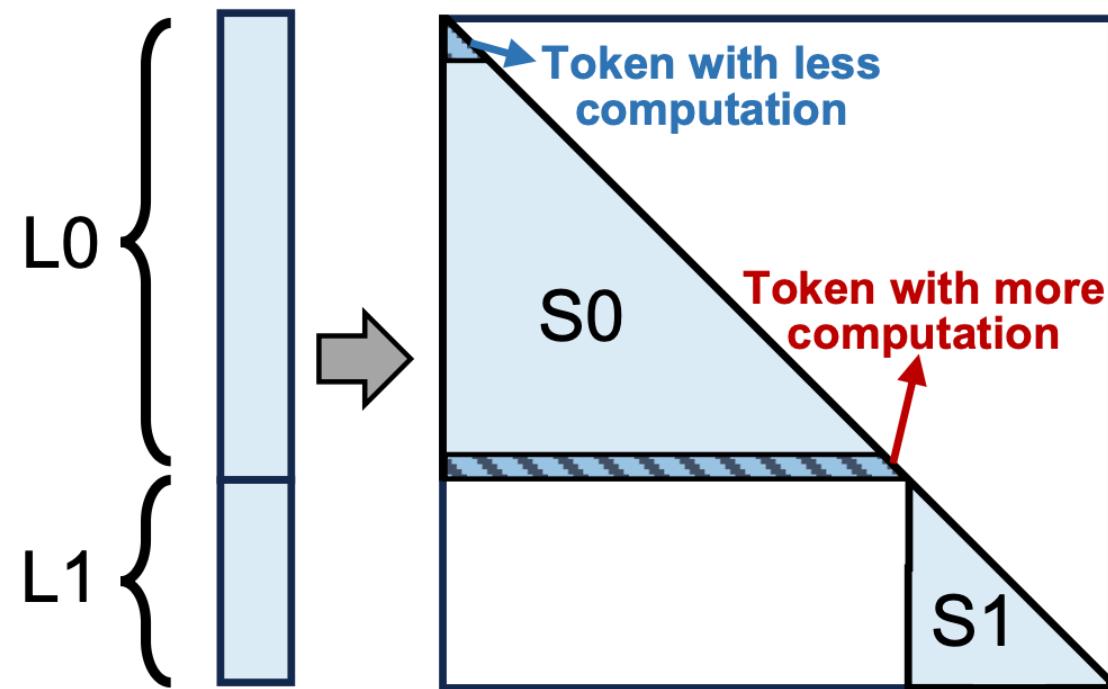
- Really big models, 405B = 910GB with 126 transformer blocks.
- Self-Attention, a layer that enriches tokens in a sequence with other token information.
- Feed-Forward, a linear matrix multiplication layer.



# Self Attention

- Attention scales quadratically with the number of past tokens  $O(T^2)$
- Not all tokens are equal: Processing later tokens in a document require more computation than earlier.
- When packing multiple documents for training, control attention span using masks.

Packed Docs      Attention Mask



# 4D Parallelism

- Huge models are trained on huge clusters because of scale and necessity.
- Tensor shape  $[B, T, H]$  across several layers.
- 4D Parallelism splits this shape in various ways.
  - Data - Split  $B$  across replicated model.
  - Pipeline - Split model into groups of layers.
  - Context - Split along sequence length  $T$
  - Tensor - Split across the  $H$ .

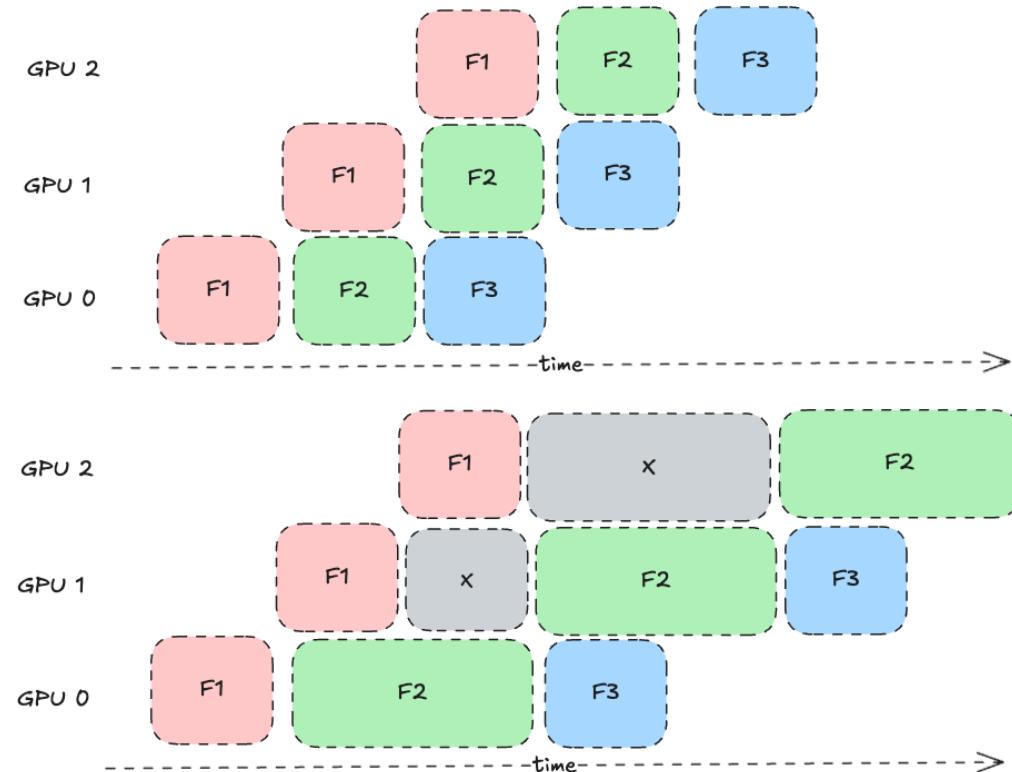
# 4D Parallelism



$DP > PP > CP > TP$

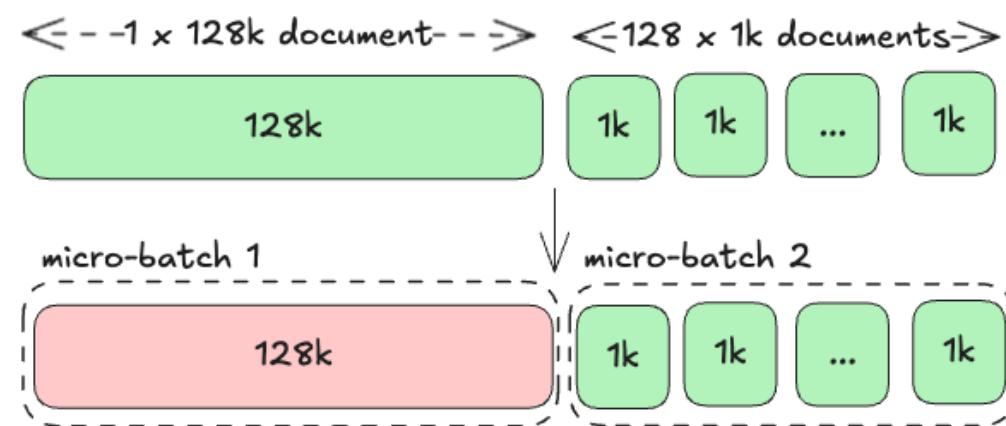
# Motivation #1 - Pipeline Imbalance

- Pipeline Parallelism splits  $B$  into  $N \mu\text{B}$  to hide delays.
- Balanced Batches  $\rightarrow$  Higher throughput



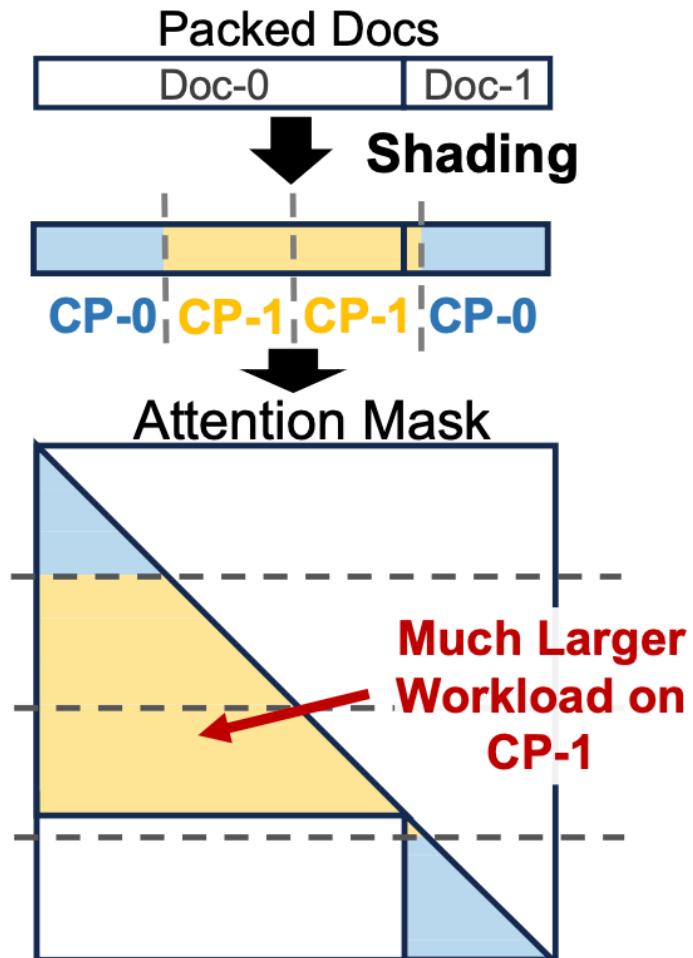
# Motivation #1 - Pipeline Imbalance

- Naive approach → split by tokens.
- Each  $\mu$ B has uniform sequence length = CONTEXT\_WINDOW\_SIZE
- Does this lead to equal attention?
- $k.128^2 \gg k.128 * 1^2$



## Motivation #2 - Context Imbalance

- Workload across CP workers should be balanced.
- After packing, split the sequence into  $2 \cdot CP$  parts and assign one from front and one from back to balance attention.
- Good heuristic, fails for multiple packed documents. Common in long context training.
- Not a huge problem, but every millisecond adds up!



- show figure 3
  - 3a motivates that there is a lot of variance in the input document size. long inputs can cause workload imbalance.
  - 3b motivates that absolute token position is actually a bad proxy for computation intensity.
    - perhaps a better metric might have been average in-document prefix length, since that is the one that matters.
    - look at their presentation, if available to see how they argue.

# Baseline: Attention-Aware packing.

**Idea:** Divide  $B$  into  $\mu B$  by estimating  $d_i^2$  as attention cost for each document.

- It works, but limited balancing improvements  
→ limited speedup.
- Higher balancing across  $\mu B$  requires balancing across multiple global  $B$ . This disturbs the random order of training and loss convergence.
- It might be impossible to come up with such a  $\mu B$  construction, if there are no candidates.

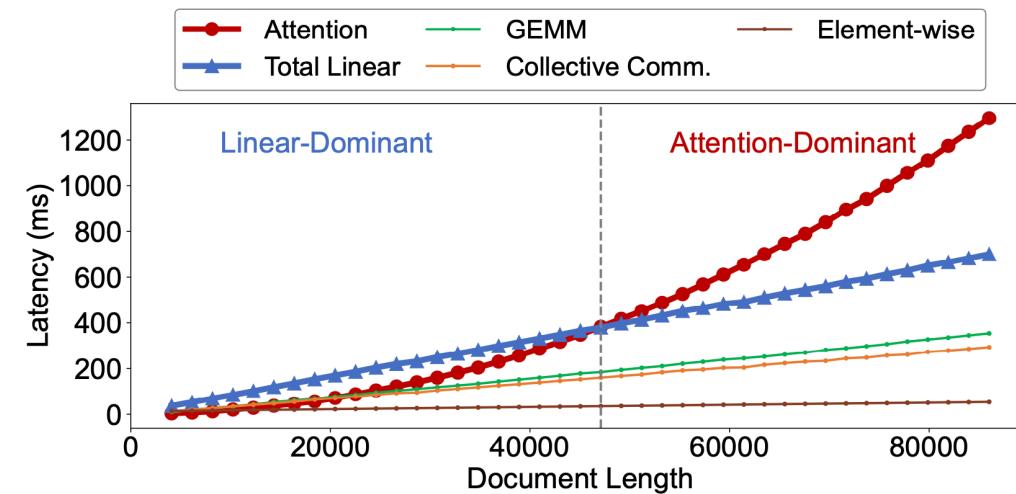
$$\begin{aligned} & \text{minimize} && \max\left(\sum_{i=1}^N x_{ij} \cdot d_i^2\right), j = 1, \dots, M \\ & \text{subject to} && \sum_{j=1}^M x_{ij} = 1, \quad i = 1, \dots, N \\ & && \sum_{i=1}^N x_{ij} \cdot d_i \leq L, \quad j = 1, \dots, M \\ & && x_{ij} \in \{0, 1\} \end{aligned}$$

# Variable-Length Packing

Idea: Allow  $\text{len}(\mu B) > \text{CONTEXT\_SIZE}$  for weaker  $\mu B$

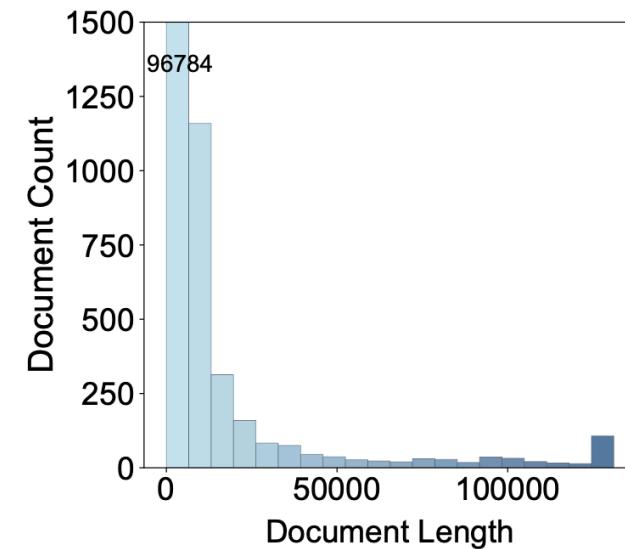
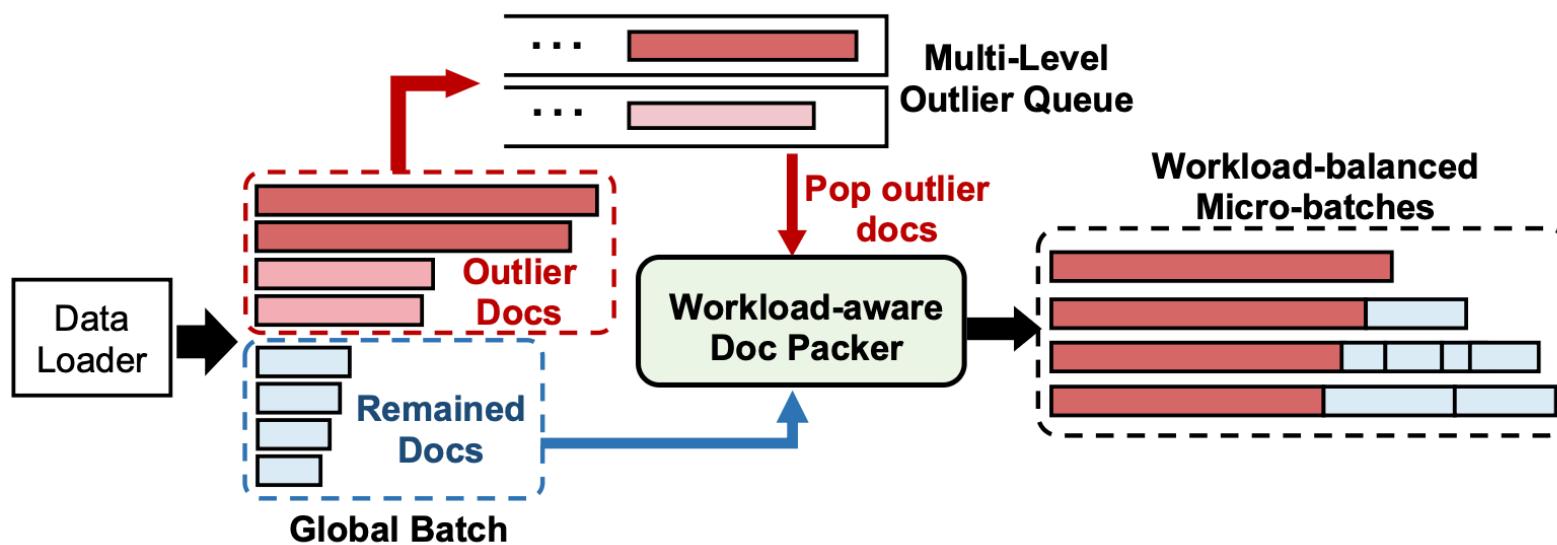
- Attention is Quadratic, but other operations are linear (feed-forward, comms etc.)
- Balance the total workload, not just attention.
- Balance long documents against many shorter documents.

$$\min \left( \max \left( \sum_{i=1}^N \left( W_a(x_{ij} \cdot d_i) + W_l(x_{ij} \cdot d_i) \right) \right) \right)$$



# Outlier document detection

- Still, you might not have sufficient smaller documents to balance the load of a long document.
- Observe: there aren't that many ultra-long documents.
- Instead of balancing across multiple batches, delay the few long documents.
- Model convergence should not hurt significantly.



# Main Algorithm

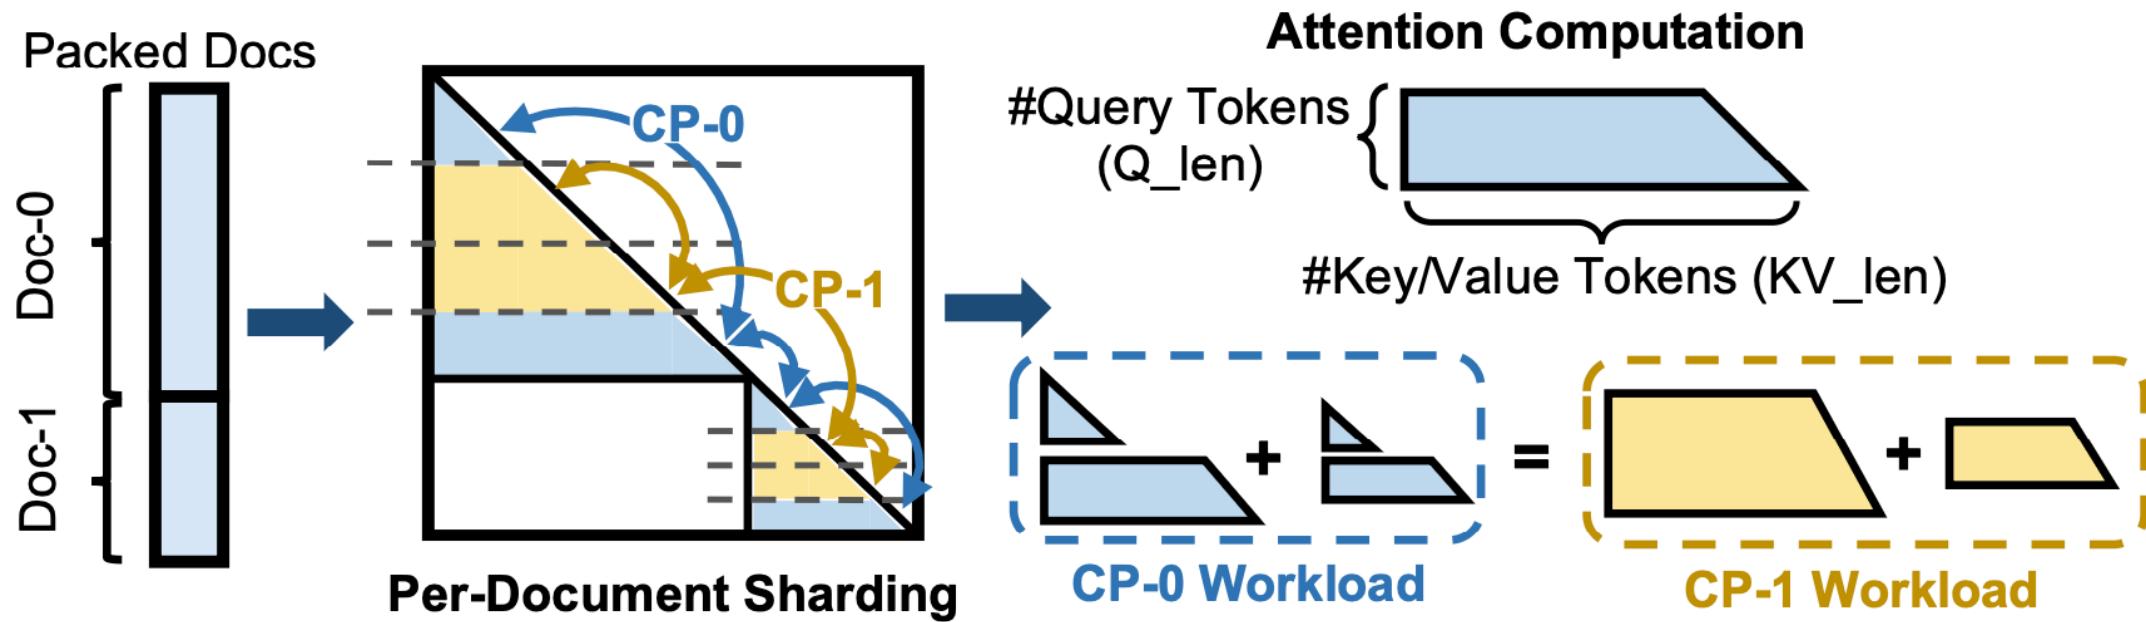
Given a DP batch

- construct a document set by removing the outlier documents, filling from the queue, and rolling over any documents from the previous iteration.
- now you effectively have a set of documents that you solved for using an ILP in S3.3, however they claim that it would take too long.
- so they instead pack greedily into an array of N micro-batches. sort the documents in reverse by length, and for each document:
  - find the least loaded microbatch (first by load, next by length)
  - if it can fit this current document, great.
  - else reserve this document for a later stage.

# Improved CP sharding

Idea: Apply CP indexing logic to each individual document.

- This should yield a more balanced workload across multiple CP workers.
- They also implement an optimization to avoid padding tokens.

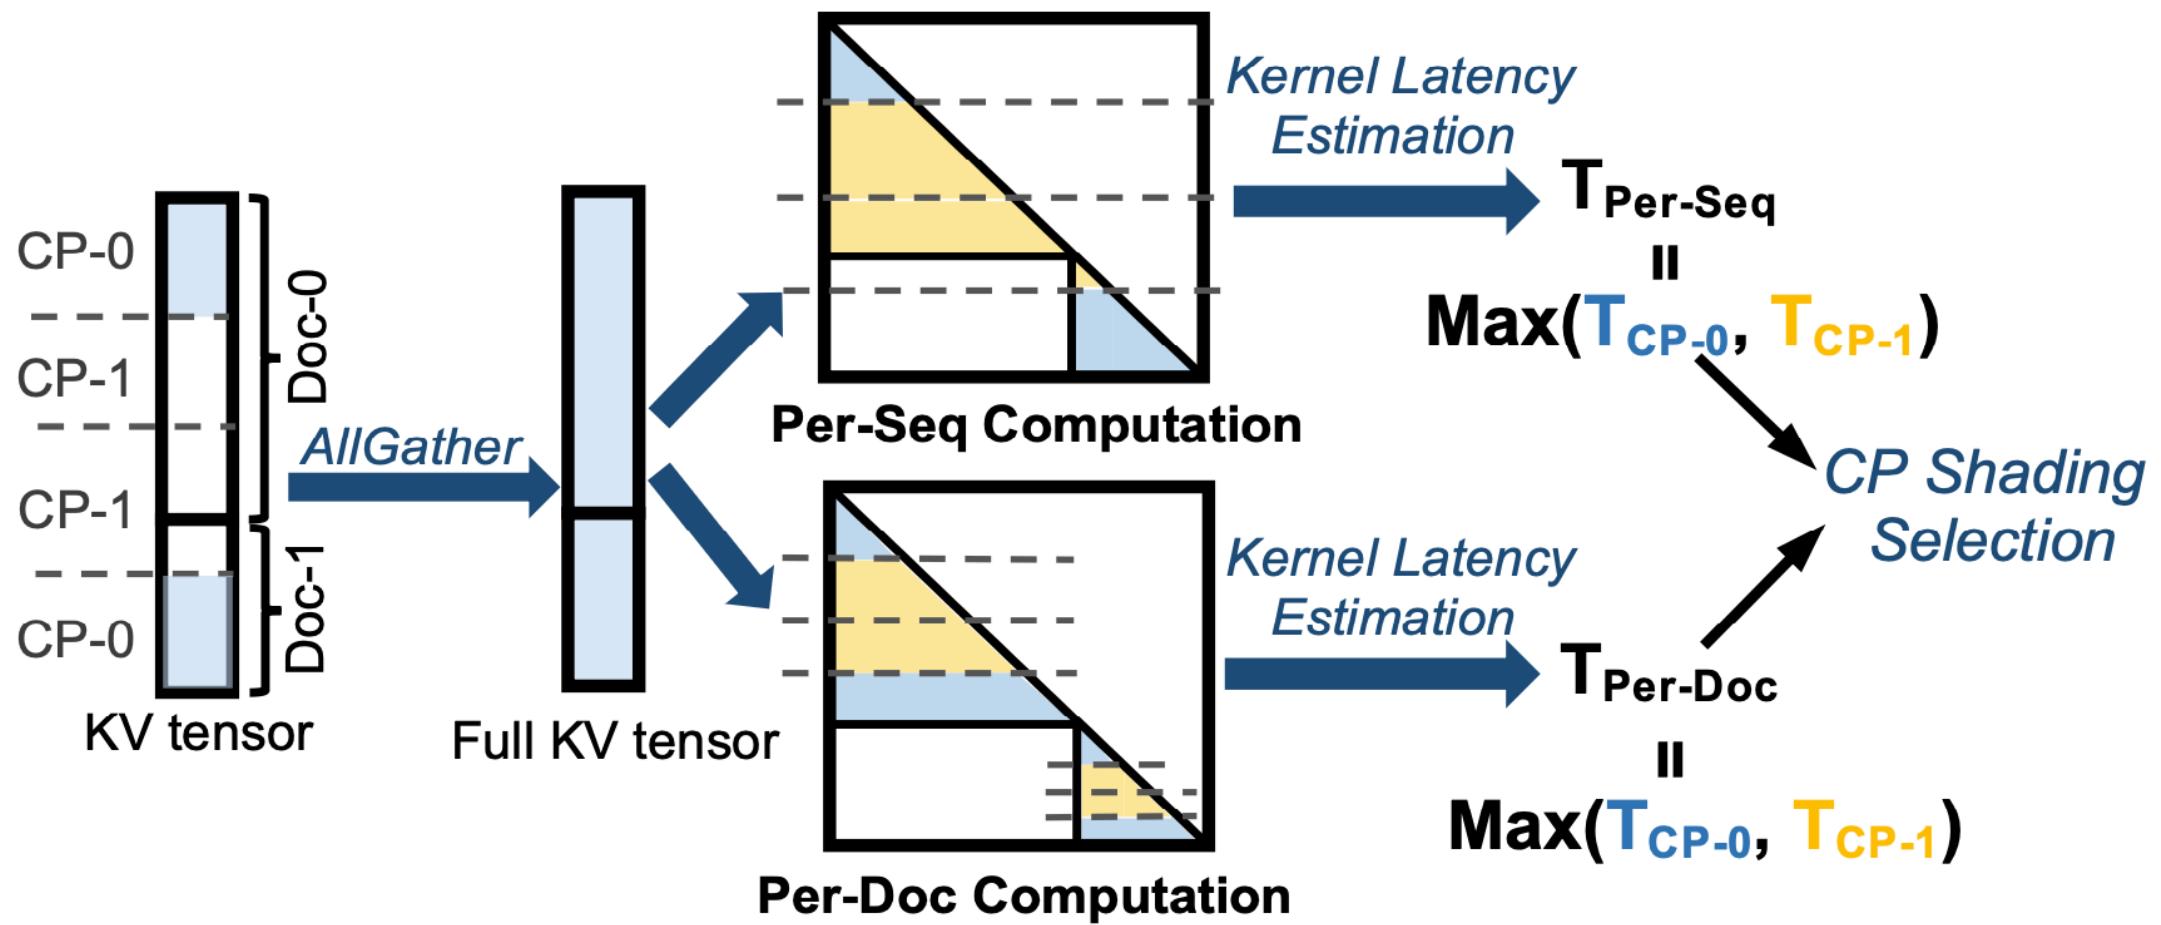


# Kernel inefficiencies

- Per-Document sharding achieves better balance, but it does not guarantee better performance.
- The attention kernel might become less efficient for smaller documents.
- Sharding lowers the  $Q$  dimension, which might cause:
  - Longer kernel time due to redundant computations (padding < 128 tokens)
  - Decreased throughput (unable to leverage TMA)

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^\top}{\sqrt{d_k}} \right) V$$

# Kernel Inefficiencies



# Evaluation

- their contributions are in PP and CP balancing. Let's think, when would PP and CP imbalance hurt most?
  - CP: longer context, more chance of long document being assigned to imbalanced worker.
  - PP: larger model, micro-batch imbalance adds more to the latency.

# Speedup Breakdown

Which optimization helps us the most?

- PP-Var-Len alone → 1.28x
- Orthogonal optimizations that combine well.
- Every second counts!

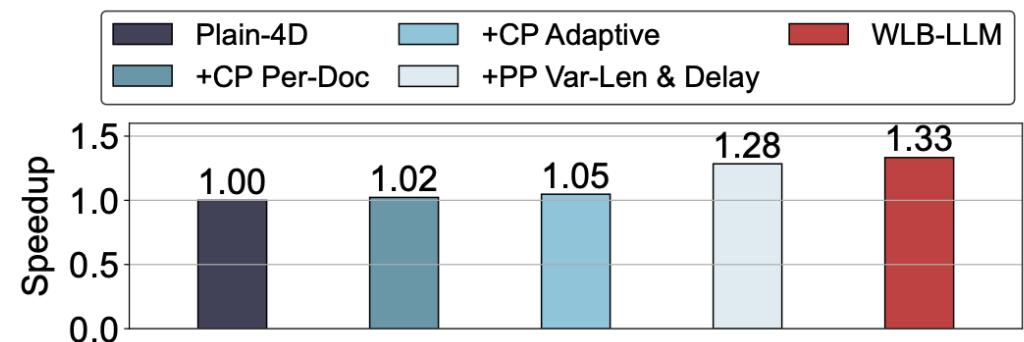


Figure 13: Performance breakdown of *WLB-LLM* on the 7B model with a 128K context window.

# Speedup across Model + Context

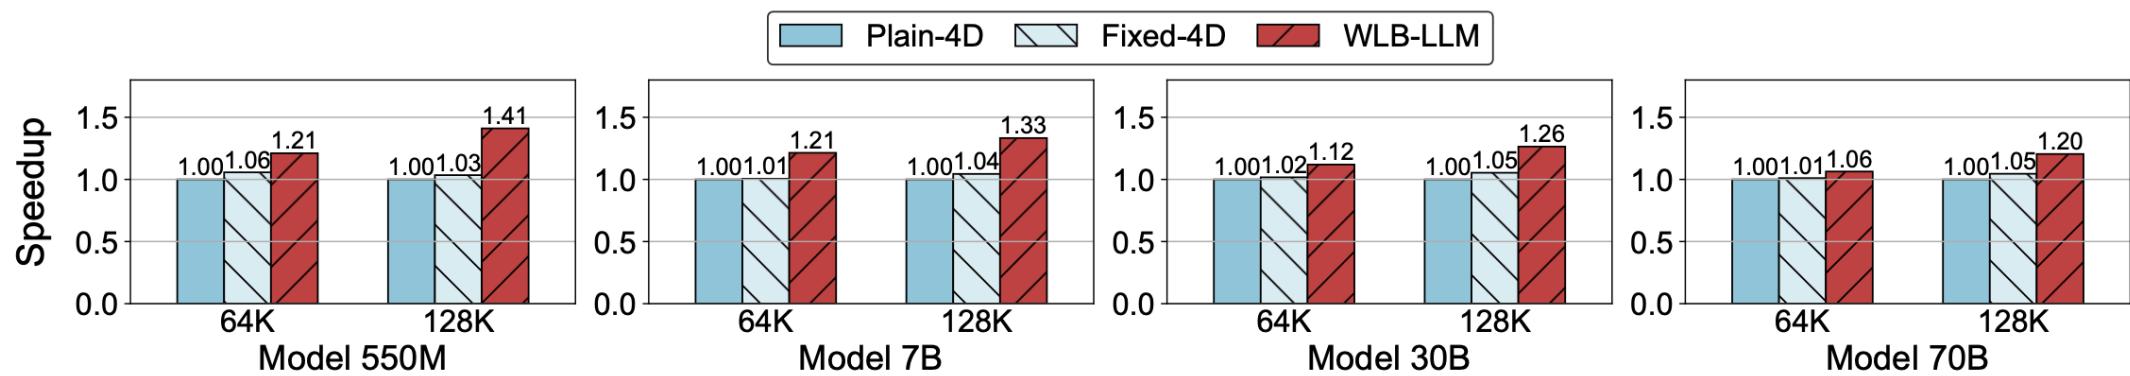


Figure 12: Training performance speedups of *WLB-LLM* and *Fixed-4D* over *Plain-4D* across various configurations.

## Speedup vs Context Window

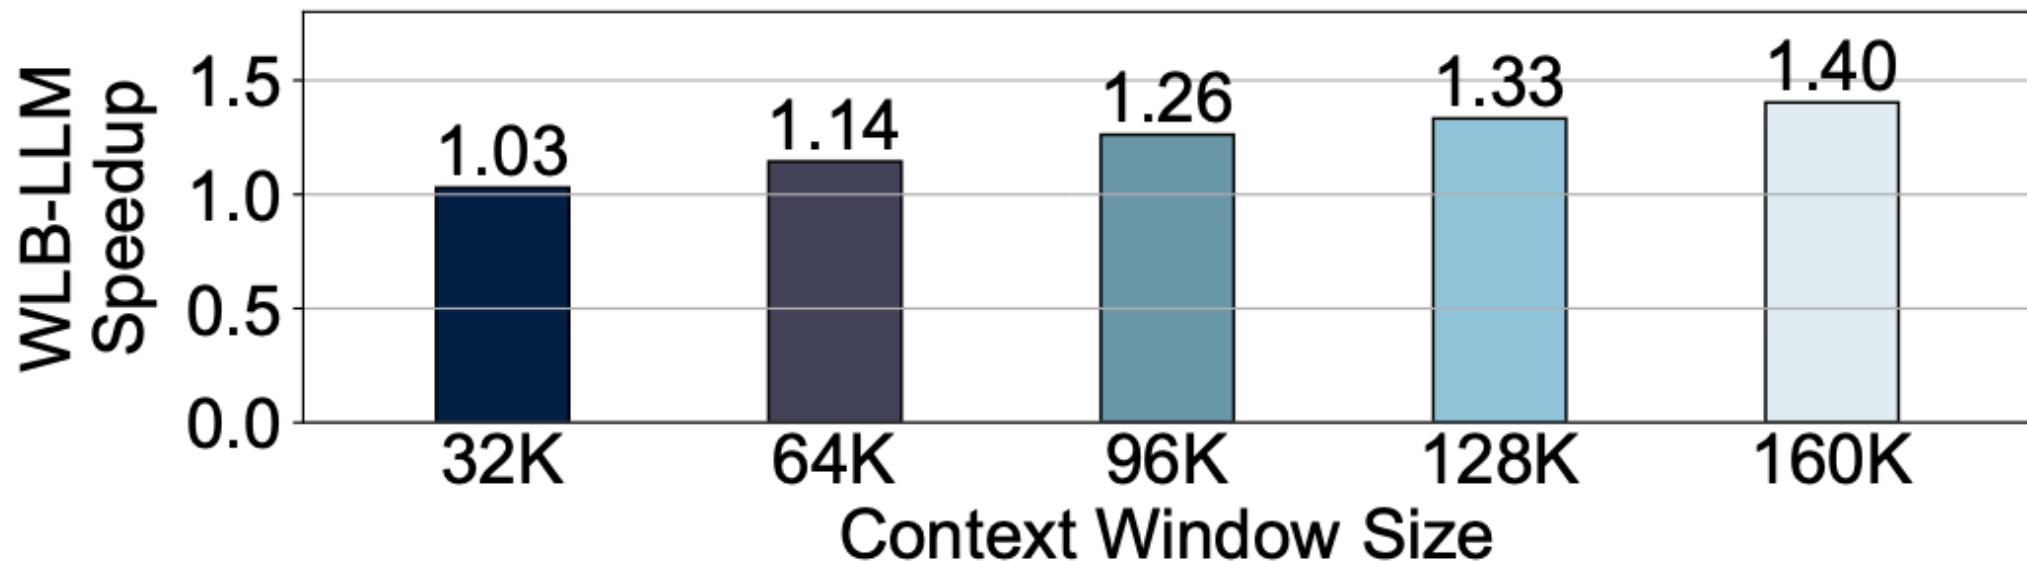


Figure 14: Speedups of *WLB-LLM* on the 7B model across context window sizes.

# Ablation - Packing Overhead

Packing Method		Imbalance Degree	Packing Overhead (ms)
Method	Config		
<i>Original Packing</i>	/	1.44	0
<i>Fixed-Len Greedy</i>	#global batch=1	1.41	4
	#global batch=2	1.22	5
	#global batch=4	1.11	5
	#global batch=8	1.08	5
<i>Fixed-Len Solver</i>	#global batch=1	1.40	467
	#global batch=2	1.18	1488
	#global batch=4	1.09	25313
<i>WLB-LLM</i>	#queue=1	1.24	8
	#queue=2	1.05	20
	#queue=3	1.05	23

Table 2: Packing imbalance degree and overhead analysis.

## Ablation - CP Sharding

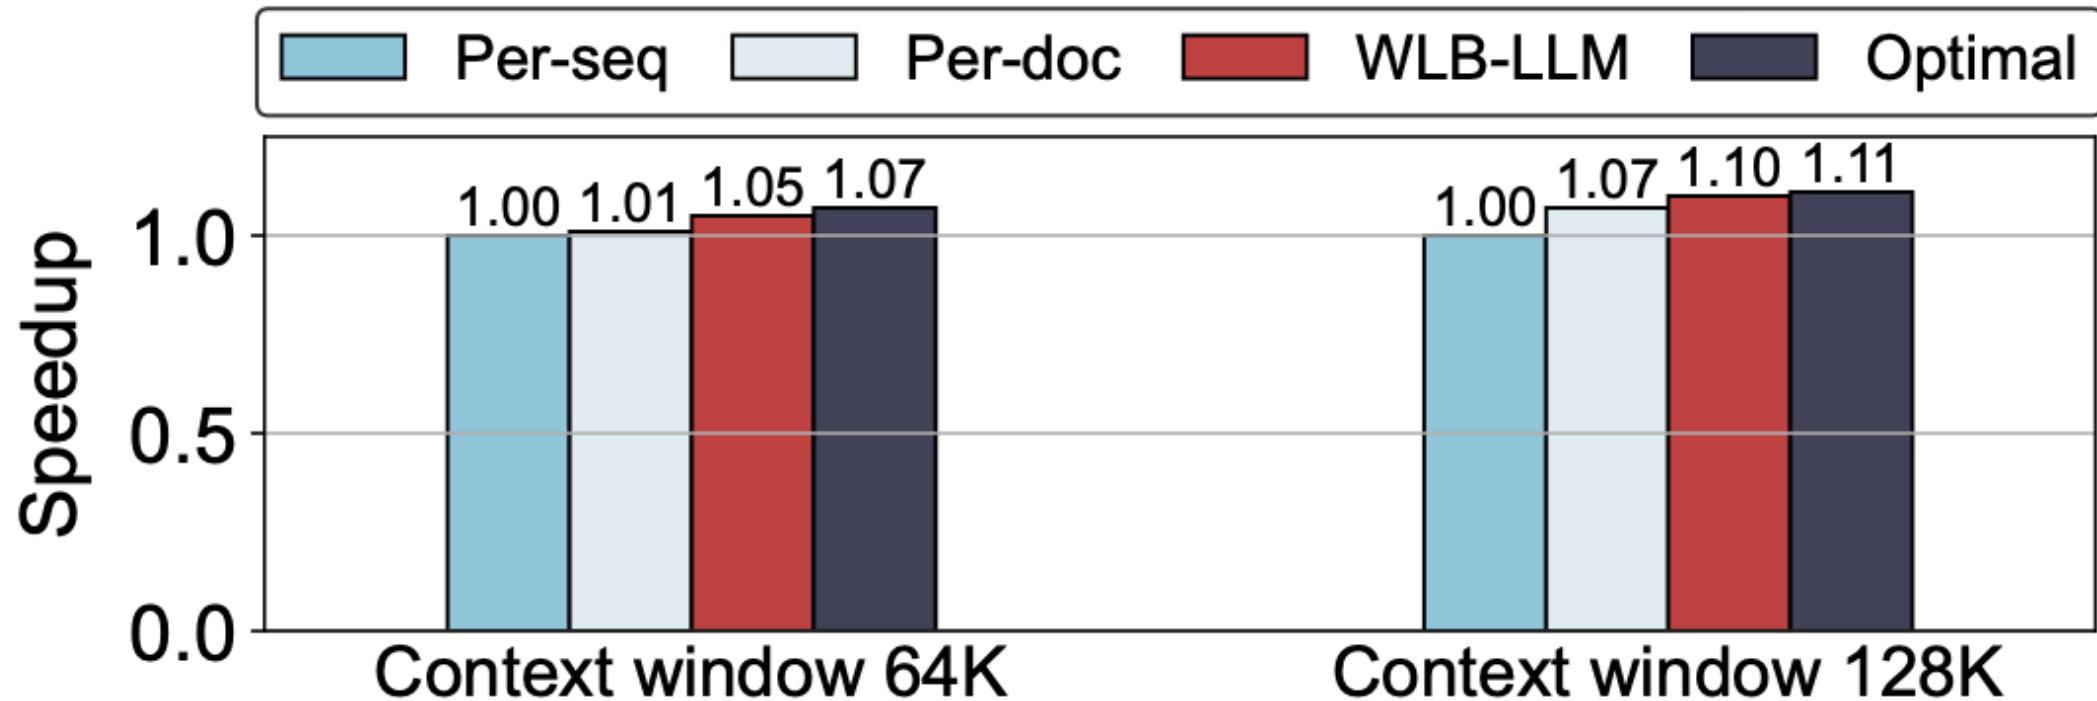


Figure 15: CP sharding performance comparison.