

# **WLB-LLM: Workload-Balanced 4D Parallelism for Large Language Model Training**

*Zheng Wang, Anna Cai, Xinfeng Xie, Zaifeng Pan, Yue Guan, Weiwei Chu, Jie Wang,  
Shikai Li, Jianyu Huang, Chris Cai, Yuchen Hao, Yufei Ding*

Kartik Ramesh

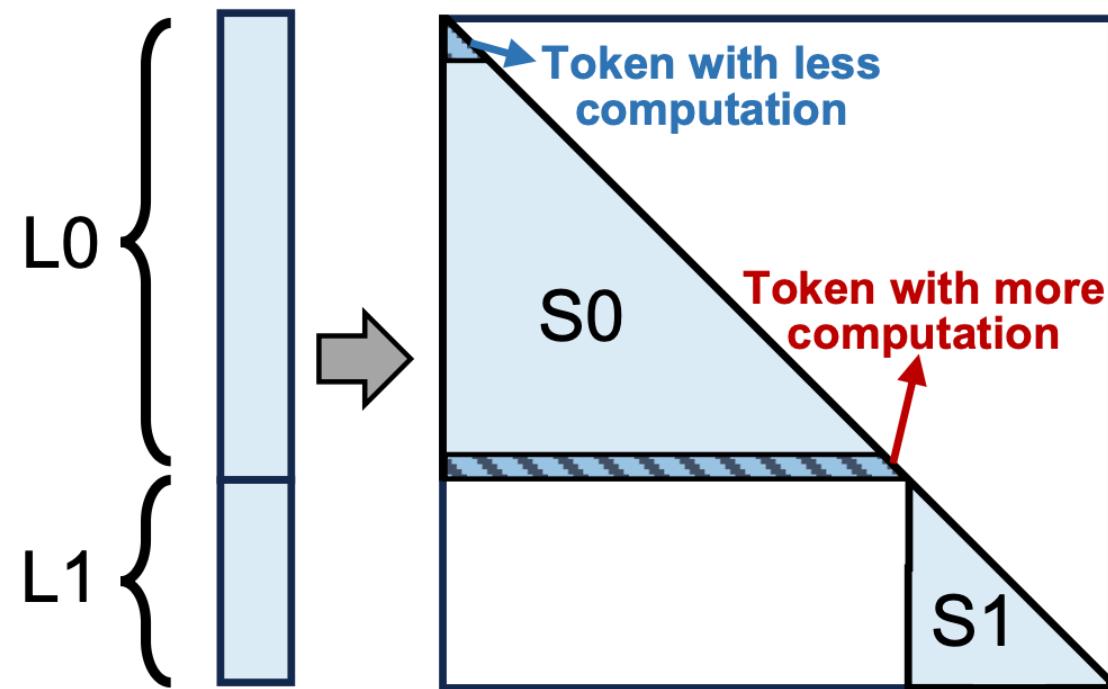
# Introduction

- Llama 3.1 405B trained over 16k GPUs for 30M H100 hours.
- Using AWS H100 pricing **\$212M**.
- A 1.2x increase in training speed would have saved **\$36M** in cloud costs.
- WLLB-LLM is a work from Meta and UCSD, released ~6 months after their training Llama 3 that achieves this.

# Self Attention

- Attention scales quadratically with the number of past tokens  $O(T^2)$
- Not all tokens are equal: Processing later tokens in a document require more computation than earlier.
- When packing multiple documents for training, control attention span using masks.

Packed Docs      Attention Mask



# 4D Parallelism

- Huge models are trained on huge clusters because of scale and necessity.
- Tensor shape  $[B, T, H]$  across several layers.
- 4D Parallelism splits this shape in various ways.
  - Data - Split  $B$  across replicated model.
  - Pipeline - Split model into groups of layers.
  - Context - Split along sequence length  $T$
  - Tensor - Split across the  $H$ .

# 4D Parallelism



$DP > PP > CP > TP$

# Key Takeaways

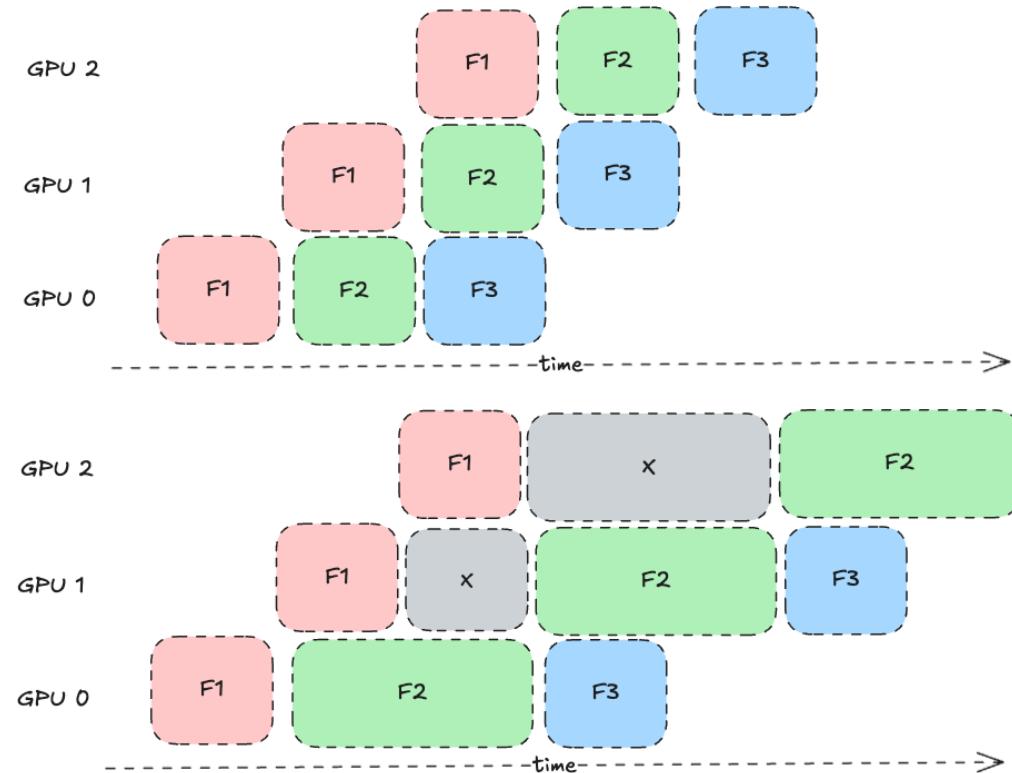
**Problem:** In long-context 4D training, token count is a weak proxy for compute; attention cost is highly non-uniform.

- **Core idea #1 (PP):** Reduce PP imbalance via attention-aware micro-batch packing.
- **Core idea #2 (CP):** Reduce CP imbalance via fine-grained per-document sharding.

**Bottom line:** WLB-LLM improves training throughput ( $\approx 1.23x$ ) without hurting convergence.

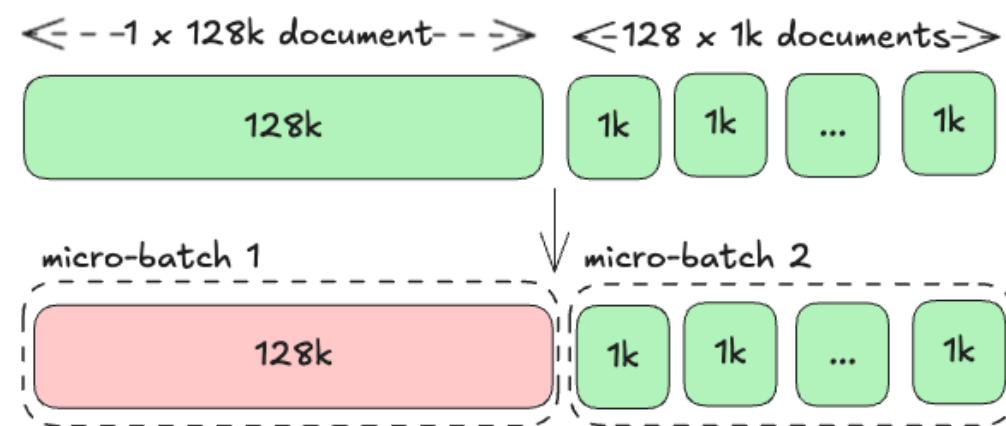
# Motivation #1 - Pipeline Imbalance

- Pipeline Parallelism splits  $B$  into  $N \mu\text{B}$  to hide delays.
- Balanced Batches  $\rightarrow$  Higher throughput



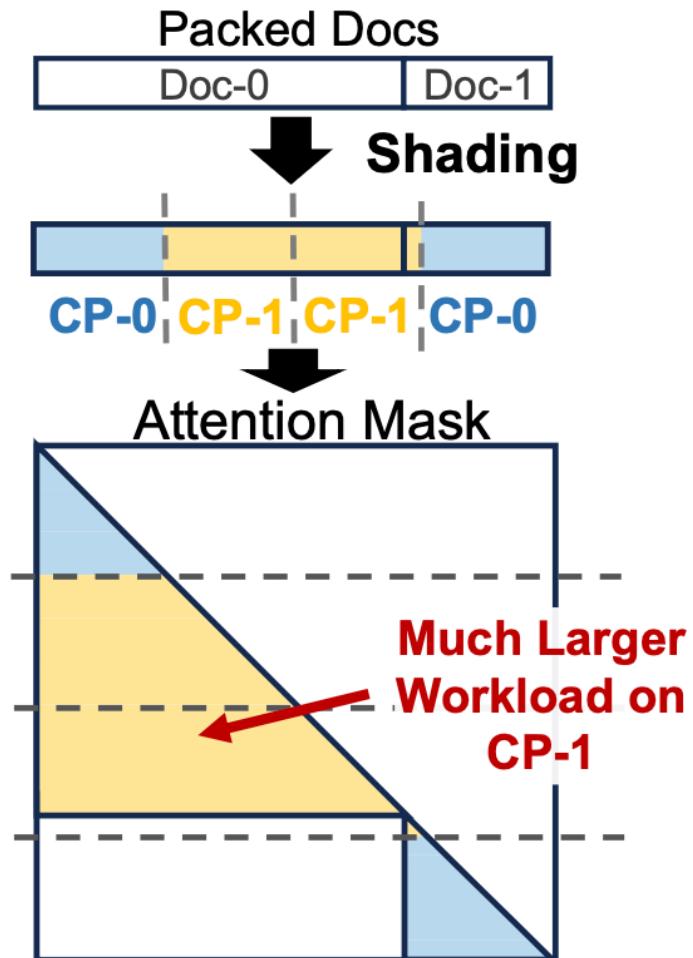
# Motivation #1 - Pipeline Imbalance

- Naive approach → split by tokens.
- Each  $\mu B$  has uniform sequence length = CONTEXT\_WINDOW\_SIZE
- Does this balance LLM workload?
- $k.128^2 \gg k.128 * 1^2$



## Motivation #2 - Context Imbalance

- Workload across CP workers should be balanced.
- After packing, split the sequence into  $2 \cdot CP$  parts and assign one from front and one from back to balance attention.
- Good heuristic, fails for multiple packed documents. Common in long context training.
- Every small delay adds up to higher-order delays.



# Baseline: Attention-Aware packing.

**Idea:** Divide  $B$  into  $\mu B$  by estimating  $d_i^2$  as attention cost for each document.

- It works, but limited balancing improvements  
→ limited speedup.
- Higher balancing across  $\mu B$  requires balancing across multiple global  $B$ . This disturbs the random order of training and loss convergence.
- It might be impossible to come up with such a  $\mu B$  construction, if there are no candidates.

$$\begin{aligned} & \text{minimize} && \max\left(\sum_{i=1}^N x_{ij} \cdot d_i^2\right), \quad j = 1, \dots, M \\ & \text{subject to} && \sum_{j=1}^M x_{ij} = 1, \quad i = 1, \dots, N \\ & && \sum_{i=1}^N x_{ij} \cdot d_i \leq L, \quad j = 1, \dots, M \\ & && x_{ij} \in \{0, 1\} \end{aligned}$$

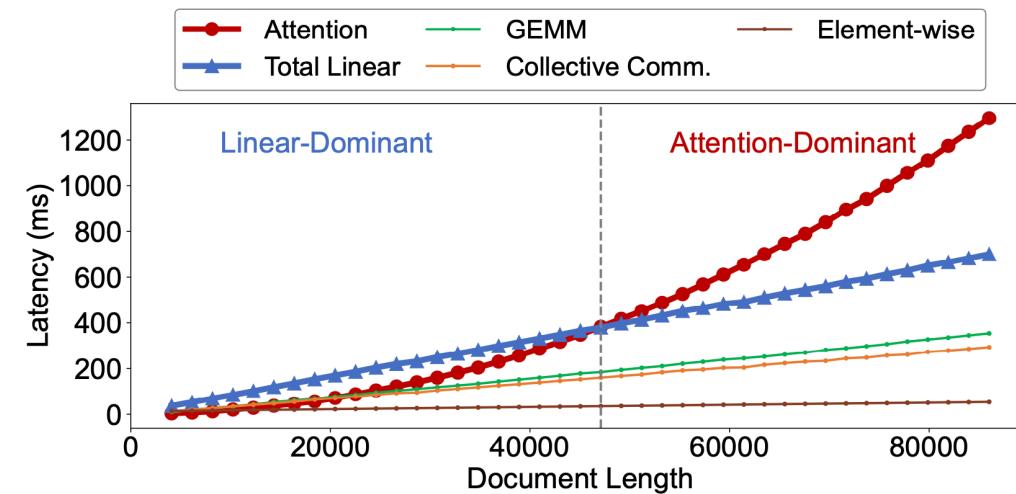
Attention Cost  
Each Document is present in only one batch  
L = Context Size

# Variable-Length Packing

Idea: Allow  $\text{len}(\mu B) > \text{CONTEXT\_SIZE}$  for weaker  $\mu B$

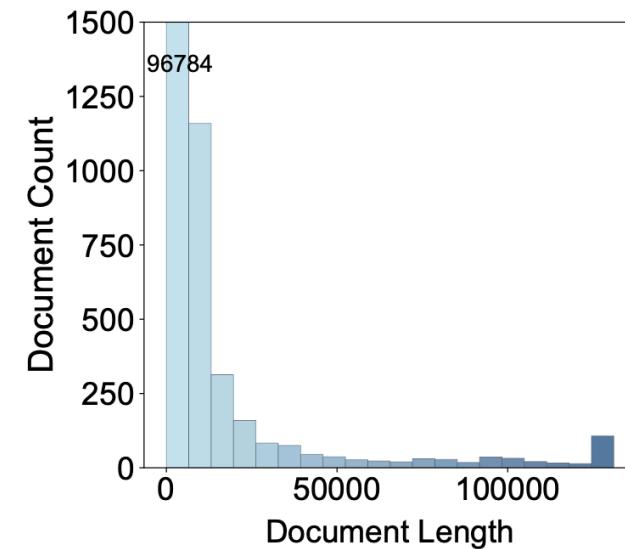
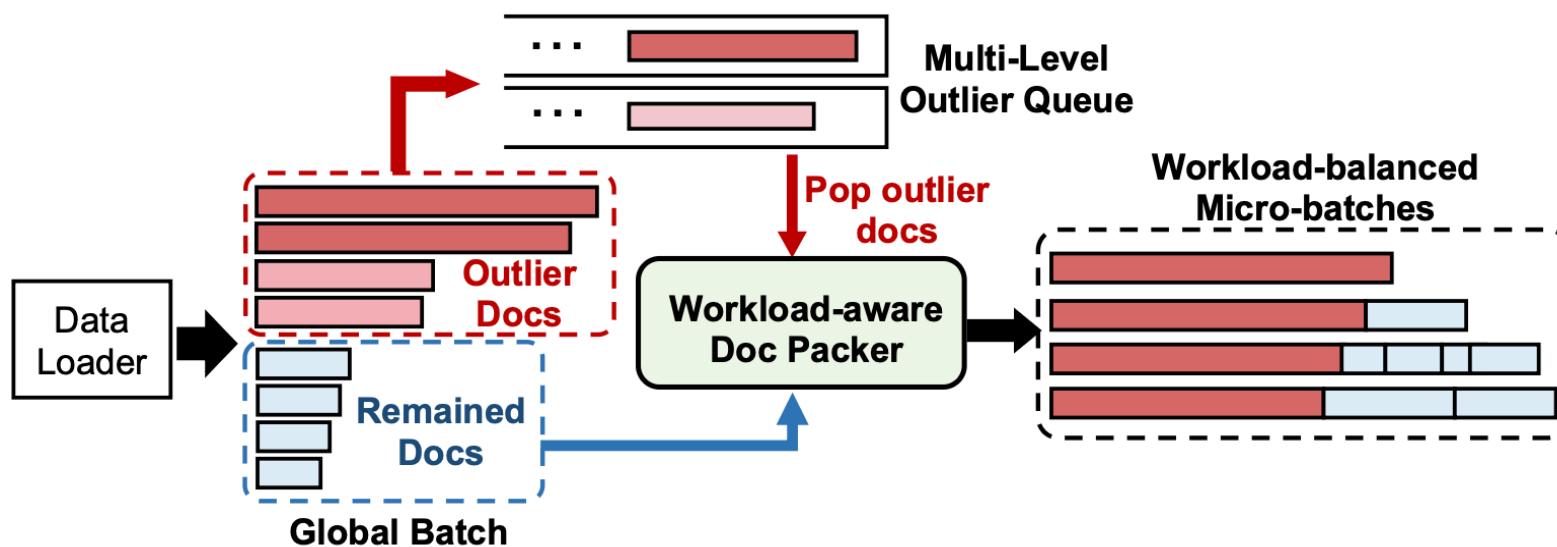
- Attention is Quadratic, but other operations are linear (feed-forward, comms etc.)
- Balance the total workload, not just attention.
- Balance long documents against many shorter documents.

$$\min \left( \max \left( \sum_{i=1}^N \left( W_a(x_{ij} \cdot d_i) + W_l(x_{ij} \cdot d_i) \right) \right) \right)$$



# Outlier document detection

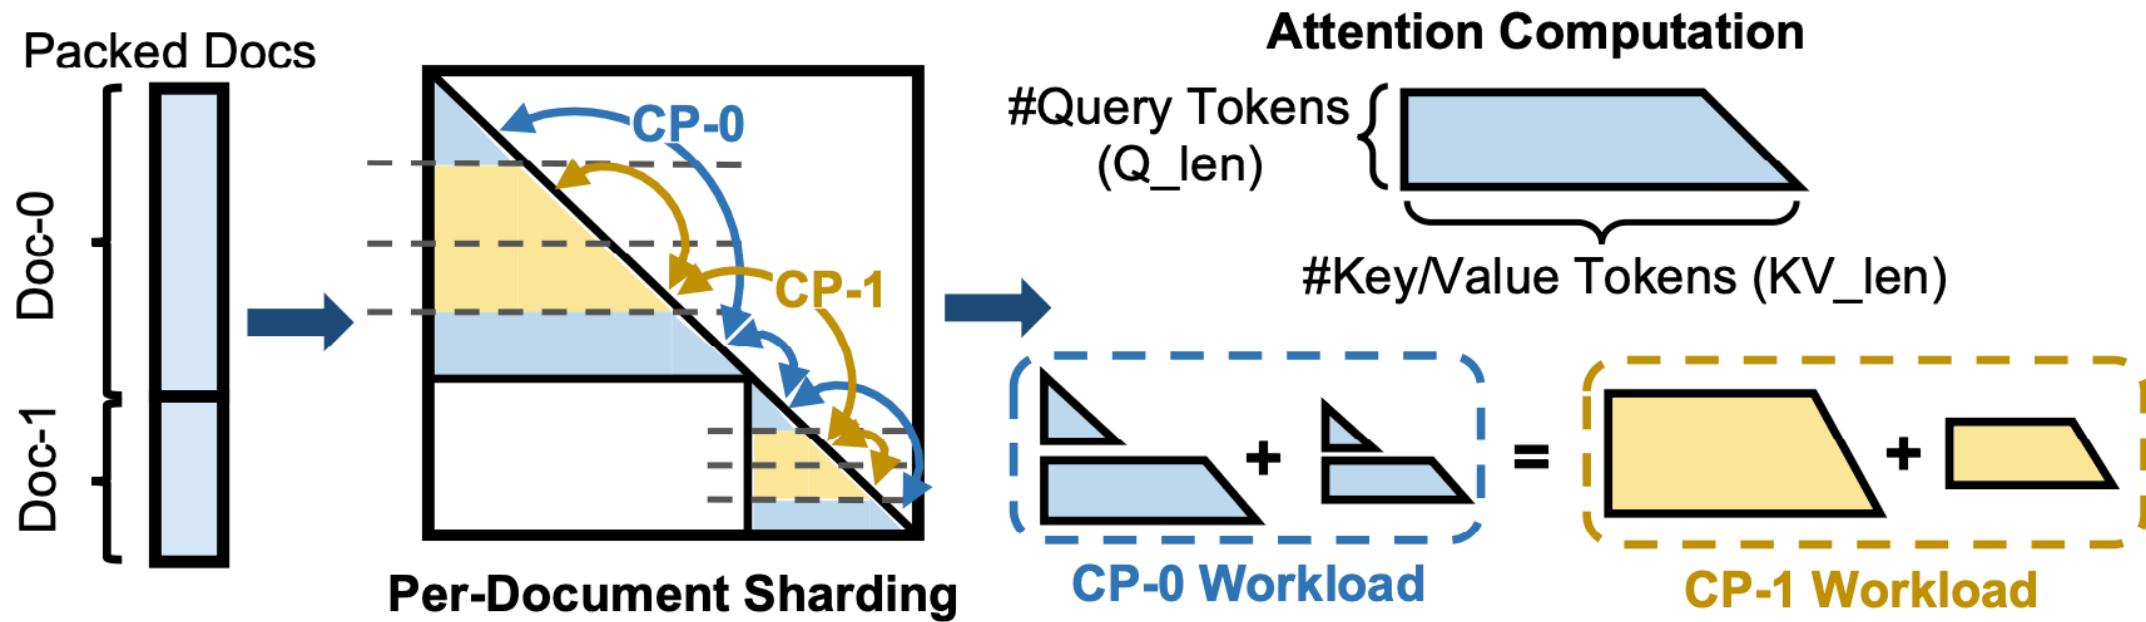
- Still, you might not have sufficient smaller documents to balance the load of a long document.
- Observe: there aren't that many ultra-long documents.
- Instead of balancing across multiple batches, delay the few long documents.
- Model convergence should not hurt significantly.



# Improved CP sharding

Idea: Apply CP indexing logic to each individual document.

- This should yield a more balanced workload across multiple CP workers.
- They also implement an optimization to avoid padding tokens.

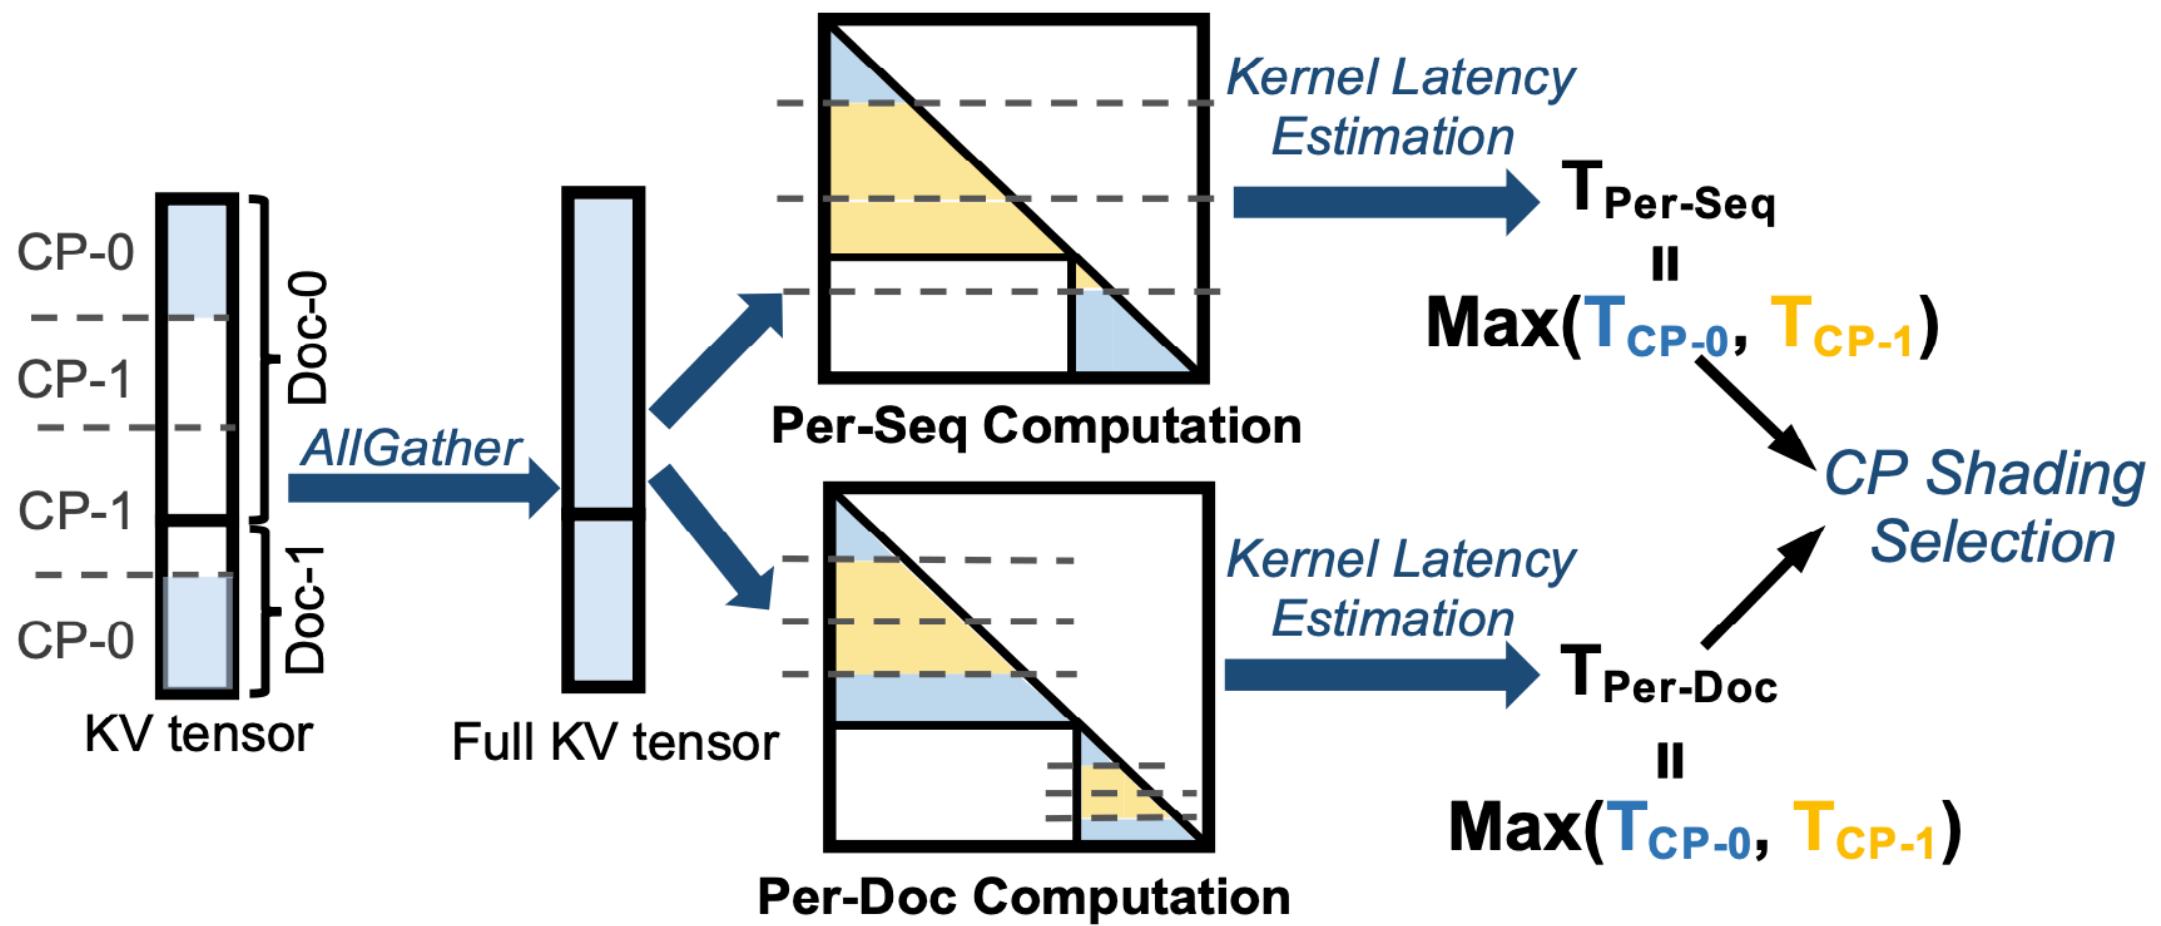


# Kernel inefficiencies

- Per-Document sharding achieves better balance, but it does not always guarantee better performance.
- Smaller per-rank attention problems reduce kernel efficiency:
  - Poor tile utilization → padding overhead for short sequences (<128 tokens).
  - Lower effective FLOPs utilization → higher time per token.
  - Reduced KV tile reuse → weaker Hopper TMA multicast benefits

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^\top}{\sqrt{d_k}} \right) V$$

# Kernel Inefficiencies



# Experimental Setup

- **Cluster:** 32 nodes, each with 8x NVIDIA H100 SXM 80GB GPUs.
- **Interconnect:** NVLink intra-node, RoCE inter-node.
- **Models:** LLaMA-like 550M, 7B, 30B, 70B; each tested at 64K and 128K context.
- **Training config:** 4D parallelism, global batch size = `PP_size x DP_size`, bfloat16 precision.
- **Baselines:**
  - `Plain-4D` : default 4D training with per-sequence CP sharding.
  - `Fixed-4D` : fixed-length packing + fixed CP sharding (per-sequence or per-document).

# Speedup Breakdown

Which optimization helps us the most?

- PP-Var-Len alone → 1.28x
- Orthogonal optimizations that combine well.
- Every second counts!

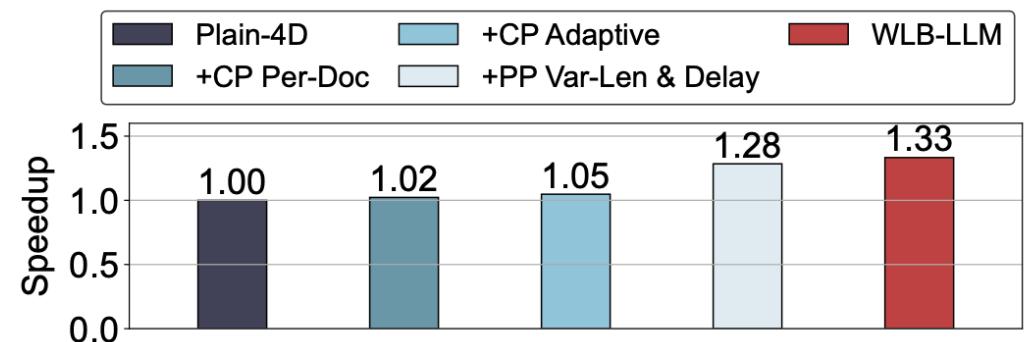


Figure 13: Performance breakdown of *WLB-LLM* on the 7B model with a 128K context window.

# Speedup across Model + Context

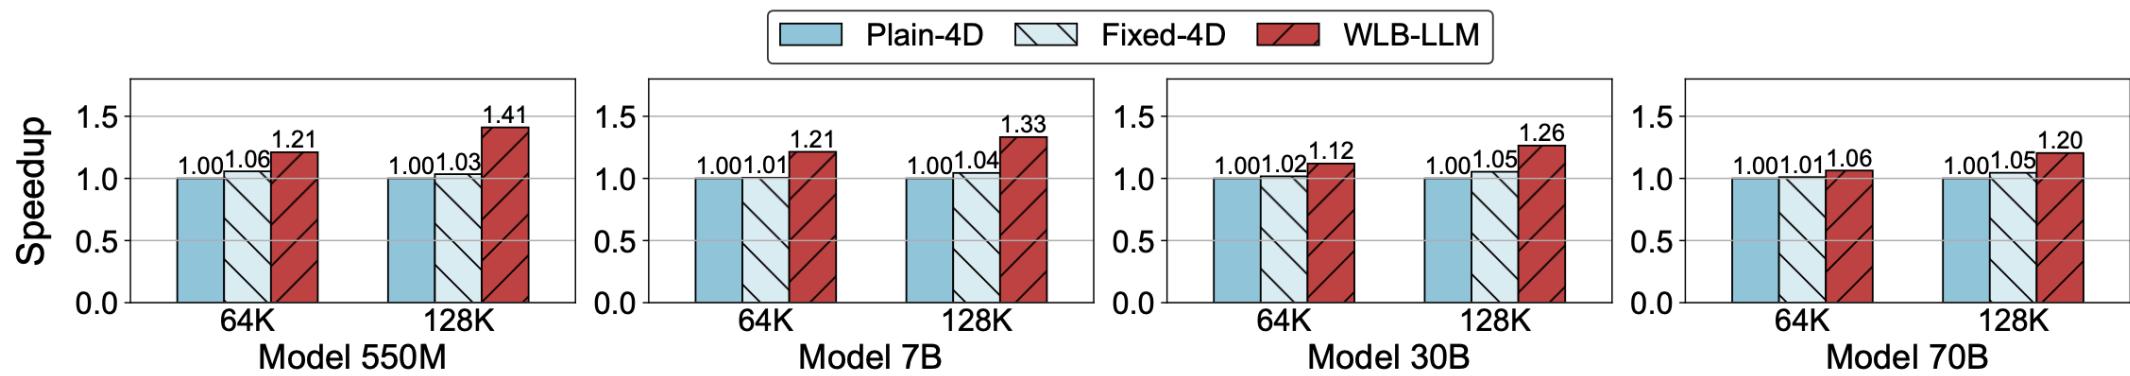


Figure 12: Training performance speedups of *WLB-LLM* and *Fixed-4D* over *Plain-4D* across various configurations.

- WLB-LLM consistently outperforms baseline for all tested configurations.
- Naive attention balancing is insufficient.
- Relative speedup decreases with increased model size.

# Other Experiments (Summary)

- **Context sensitivity (Fig. 14):** Speedup increases with context length (about  $1.07x$  @64K to  $1.40x$  @160K on 7B), consistent with worse imbalance at longer contexts.
- **Packing overhead vs balance (Table 2):** WLB-LLM reaches near-optimal imbalance with low runtime overhead (tens of ms), while solver-based packing can be prohibitively slow.
- **CP sharding ablation (Fig. 15):** Adaptive sharding consistently outperforms always-per-sequence or always-per-document sharding.
- **Convergence / quality:** Their training-loss curves indicate no clear quality regression from the system optimizations.

# Discussion & Critique

## Strengths

- Identifies and fixes a bottleneck for training long-context Llama models with 4D parallelism.
- Joint PP+CP optimization gives meaningful real-system gains, especially as context windows grow.
- Engineering is practical: low overhead and no obvious convergence regression in their reported runs.

## Weaknesses

- Workload dependence and limited generalization evidence outside their evaluated distribution.
- Strong dependence on a small number of extreme outliers; unclear benefit when length distributions are flatter.
- Heavy use of heuristics (packing + sharding selection) without strong guarantees in worst-case settings.

# Related Work

## (1) Efficient Long-context Language Model Training by Core Attention Disaggregation (DistCA)

- Split out “core attention” ( $\text{softmax}(QK^\top)V$ ) as a weightless compute service, separate from the rest of the transformer.
- Better than WLB-style baselines at scale: reports  $\sim 1.15\text{--}1.35\times$  throughput gains over their WLB “ideal” baseline in 4D (with PP), depending on workload.

## (2) ByteScale Efficient Scaling of LLM Training with a 2048K Context Length on More Than 12,000 GPUs

- Hybrid Data Parallelism (HDP): unify DP + CP into one dynamic device mesh.
- Length-aware sharding: use the minimum number of devices per sequence.
- Short sequences stay local (skip CP comm), long sequences shard across more GPUs.

## (3) Ordering efficiency

- Reduce pipeline bubbles by optimizing scheduling.
- PipeDream, 1F1B, Seq1F1B.

# What did you think?

**Problem:** In long-context 4D training, token count is a weak proxy for compute; attention cost is highly non-uniform.

- **Core idea #1 (PP):** Reduce PP imbalance via attention-aware micro-batch packing.
- **Core idea #2 (CP):** Reduce CP imbalance via fine-grained per-document sharding.

**Bottom line:** WLB-LLM improves training throughput ( $\approx 1.23x$ ) without hurting convergence.

Packing Method		Imbalance Degree	Packing Overhead (ms)
Method	Config		
<i>Original Packing</i>	/	1.44	0
<i>Fixed-Len Greedy</i>	#global batch=1	1.41	4
	#global batch=2	1.22	5
	#global batch=4	1.11	5
	#global batch=8	1.08	5
<i>Fixed-Len Solver</i>	#global batch=1	1.40	467
	#global batch=2	1.18	1488
	#global batch=4	1.09	25313
<i>WLB-LLM</i>	#queue=1	1.24	8
	#queue=2	1.05	20
	#queue=3	1.05	23

Table 2: Packing imbalance degree and overhead analysis.